

music

April 1, 2025

1 Set Up

```
[198]: # Import
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from xgboost import XGBClassifier
from sklearn.ensemble import IsolationForest
from sklearn.svm import SVC # Support Vector Classification
import lightgbm as lgb
from imblearn.ensemble import BalancedRandomForestClassifier
from sklearn.metrics import classification_report, f1_score, recall_score, \
    ↪confusion_matrix, ConfusionMatrixDisplay
from imblearn.over_sampling import SMOTE
from sklearn.impute import SimpleImputer
```

```
[199]: # Functions
def most_frequent(series):
    return series.mode()[0] if not series.mode().empty else None
```

```
[200]: # Settings

pd.set_option('display.max_columns', None) # Show all columns
pd.set_option('display.width', 1000) # Prevent line breaks
pd.set_option('display.max_colwidth', None) # Show full column values
```

```
[204]: # Read Files

# Path to your Parquet file
parquet_file = '/Users/henryding/Documents/Seminar/Data/output_data1.parquet'

# Read the Parquet file into a DataFrame
```

```
df = pd.read_parquet(parquet_file, engine='pyarrow') # You can also use
↳ 'fastparquet'

# Display the first few rows of the DataFrame
display(df.head(1000))
num_rows = df.shape[0]
print(f"The number of rows in the DataFrame is: {num_rows}")
```

	status	gender	length	firstName	level	lastName	registration	userId
	ts	auth	page	sessionId				
	location	itemInSession						
	userAgent							
						song		
						artist	method	
0	200	M	524.32934	Shlok	paid	Johnson	1.533735e+12	1749042
	1538352001000	Logged In	NextSong	22683		Dallas-Fort		
	Worth-Arlington, TX		278			"Mozilla/5.0		
	(Windows NT 6.1; WOW64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/36.0.							
	1985.143 Safari/537.36"							
						Ich mache einen Spiegel -		
	Dream Part 4							
	Popol Vuh	PUT						
1	200	F	238.39302	Vianney	paid	Miller	1.537500e+12	1563081
	1538352002000	Logged In	NextSong	20836		San		
	Francisco-Oakland-Hayward, CA		9			"Mozilla/5.0		
	(Macintosh; Intel Mac OS X 10_9_4) AppleWebKit/537.36 (KHTML, like Gecko)							
	Chrome/36.0.1985.143 Safari/537.36"							
	MiÃ	ntele						
	Los Bunkers	PUT						
2	200	F	140.35546	Vina	paid	Bailey	1.536415e+12	1697168
	1538352002000	Logged In	NextSong	4593				
	Hilo, HI		109					
	Mozilla/5.0 (Macintosh; Intel Mac OS X 10.9; rv:31.0) Gecko/20100101							
	Firefox/31.0							
						Baby Talk		
						Lush	PUT	
3	200	M	277.15873	Andres	paid	Foley	1.534387e+12	1222580
	1538352003000	Logged In	NextSong	6370				
	Watertown, SD		71			"Mozilla/5.0 (Macintosh; Intel		
	Mac OS X 10_9_4) AppleWebKit/537.77.4 (KHTML, like Gecko) Version/7.0.5 Safari/							
	537.77.4"							
						Horn Concerto No. 4 in E flat K495: II. Romance (Andante cantabile)		
	Barry Tuckwell/Academy of St Martin-in-the-Fields/Sir Neville Marriner	PUT						

```

4      200      F  1121.25342  Aaliyah  paid  Ramirez  1.537381e+12  1714398
↳ 1538352003000  Logged In      NextSong      22316
↳ Baltimore-Columbia-Towson, MD      21
↳ "Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit/537.36 (KHTML, like Gecko)
↳ Chrome/36.0.1985.143 Safari/537.36"  Close To The Edge (I. The Solid Time Of
↳ Change_ II. Total Mass Retain_ III. I Get Up I Get Down_ IV. Seasons Of Man)
↳ (Remastered LP Version)
↳      Yes      PUT
..      ...      ...      ...      ...      ...      ...      ...      ...      ...
↳      ...      ...      ...
↳
↳
↳
↳
↳
↳
↳
↳
995      307      F      NaN      Alivia  paid  Williams  1.535955e+12  1792538
↳ 1538352525000  Logged In      Thumbs Up      8871  New York-Newark-Jersey
↳ City, NY-NJ-PA      75
↳      Mozilla/5.0 (Macintosh; Intel Mac OS X 10.9; rv:31.0) Gecko/20100101
↳ Firefox/31.0
↳
↳
↳      None
↳      None      PUT
996      200      M  315.48036  Christian  free      Klein  1.536940e+12  1291813
↳ 1538352526000  Logged In      NextSong      14821  New York-Newark-Jersey
↳ City, NY-NJ-PA      10      "Mozilla/5.0 (Macintosh;
↳ Intel Mac OS X 10_9_2) AppleWebKit/537.74.9 (KHTML, like Gecko) Version/7.0.2
↳ Safari/537.74.9"
↳
↳
↳ Odessa      Caribou
↳      PUT
997      200      M  248.73751  Kristofer  free      James  1.536879e+12  1929921
↳ 1538352526000  Logged In      NextSong      9831
↳ Seattle-Tacoma-Bellevue, WA      56
↳      Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:31.0)
↳ Gecko/20100101 Firefox/31.0
↳
↳      Jigsaw
↳ Falling Into Place
↳ Radiohead      PUT
998      200      F  53.08036      Zoey  free      Gregory  1.533190e+12  1839740
↳ 1538352527000  Logged In      NextSong      24848
↳ Phoenix-Mesa-Scottsdale, AZ      11  "Mozilla/5.0 (X11; Linux x86_64)
↳ AppleWebKit/537.36 (KHTML, like Gecko) Ubuntu Chromium/34.0.1847.116 Chrome/34.
↳ 0.1847.116 Safari/537.36"
↳
↳      Broken Hearted Hoover Fixer
↳ Sucker Guy      Glen
↳ Hansard      PUT

```

```

999      200      F      NaN      Zoey free Gregory 1.533190e+12 1839740
↳ 1538352527000 Logged In Roll Advert      24848
↳ Phoenix-Mesa-Scottsdale, AZ      12 "Mozilla/5.0 (X11; Linux x86_64)
↳ AppleWebKit/537.36 (KHTML, like Gecko) Ubuntu Chromium/34.0.1847.116 Chrome/34.
↳ 0.1847.116 Safari/537.36"
↳
↳      None
↳ None      GET

```

[1000 rows x 18 columns]

The number of rows in the DataFrame is: 26259199

2 Exploratory Analysis

[205]: *# Categorical Variables*

```

unique_level = df['level'].unique()
print(f"level: {unique_level}")

unique_auth = df['auth'].unique()
print(f"auth: {unique_auth}")

unique_method = df['method'].unique()
print(f"method: {unique_method}")

unique_page = df['page'].unique()
print(f"page: {unique_page}")

```

```

level: ['paid' 'free']
auth: ['Logged In' 'Logged Out' 'Cancelled' 'Guest']
method: ['PUT' 'GET']
page: ['NextSong' 'Home' 'Thumbs Up' 'Error' 'Help' 'Add to Playlist'
'Roll Advert' 'Login' 'Thumbs Down' 'Settings' 'About' 'Add Friend'
'Upgrade' 'Submit Upgrade' 'Logout' 'Downgrade' 'Save Settings'
'Submit Downgrade' 'Cancel' 'Cancellation Confirmation' 'Register'
'Submit Registration']

```

[206]: *# User Usage*

```

user_counts = df.groupby('userId').size().reset_index(name='count')
display(user_counts.sort_values(by='count', ascending=False))

```

	userId	count
5936	1261737	778479
12534	1564221	13591
20802	1931933	12831
163	1006695	12372
7562	1336969	11858

...	...	...
6072	1267517	1
9698	1434698	1
17778	1793623	1
9129	1408726	1
15322	1689121	1

[22278 rows x 2 columns]

```
[207]: # Time Period
# Convert Timestamps to datetime
df['ts'] = pd.to_datetime(df['ts'], unit='ms')

print(f"timestamp min: {df['ts'].min()}")
print(f"timestamp max: {df['ts'].max()}")


# Convert Registration to datetime
df['registration'] = pd.to_datetime(df['registration'], unit='ms')

print(f"registration min: {df['registration'].min()}")
print(f"registration max: {df['registration'].max()}")
```

```
timestamp min: 2018-10-01 00:00:01
timestamp max: 2018-12-01 00:00:02
registration min: 2017-10-14 22:05:25
registration max: 2018-12-03 07:23:42
```


```
[208]: # User Registration
grouped_registration = df.groupby('userId')['registration'].agg('min')
display(grouped_registration.head())
plt.hist(grouped_registration, bins=100)
plt.title('User Registration')
plt.show()
```

```
userId
1000025    2018-07-10 09:30:08
1000035    2018-09-12 19:28:22
1000083    2018-09-07 18:01:49
1000103    2018-09-22 07:27:25
1000164    2018-08-12 09:32:01
Name: registration, dtype: datetime64[ns]
```




```
[11]: # First Activity Date
grouped_start = df.groupby('userId')['ts'].agg('min')
display(grouped_start.head())
plt.hist(grouped_start, bins=100)
plt.xticks(rotation=45)
plt.title('First Activity Date')
plt.show()
```

```
userId
1000025    2018-10-02 08:59:29
1000035    2018-10-05 18:29:46
1000083    2018-10-01 05:37:45
1000103    2018-10-04 17:31:17
1000164    2018-10-01 17:40:18
Name: ts, dtype: datetime64[ns]
```



```
[12]: # Last Activity Date
grouped_end = df.groupby('userId')['ts'].agg('max')
display(grouped_end.head())
plt.hist(grouped_end, bins=100)
plt.xticks(rotation=45)
plt.title('Last Activity Date')
plt.show()
```

```
userId
1000025    2018-10-18 20:33:05
1000035    2018-11-20 09:14:13
1000083    2018-10-12 10:04:58
1000103    2018-11-21 03:01:43
1000164    2018-11-30 16:53:06
Name: ts, dtype: datetime64[ns]
```



```
[101]: # High-Cardinality Variables
song_count = df.groupby('song').size().count()
print(f"Number of Songs: {song_count}")

artist_count = df.groupby('artist').size().count()
print(f"Number of Artists: {artist_count}")

user_count = df.groupby('userId').size().count()
print(f"Number of Users: {user_count}")
```


```
Number of Songs: 253564
Number of Artists: 38337
Number of Users: 22278
```

```
[13]: # Continuous Variables

plt.hist(df['itemInSession'], bins=100)
plt.title("itemInSession")
```



```
plt.show()
```



3 Data Preprocessing

```
[ ]:
```

4 Feature Engineering

```
[209]: # Create "week" feature to group by week
df['week'] = df['ts'].dt.to_period('W').astype('datetime64[ns]')

# Sort by userId and timestamp to calculate session_diff properly
df = df.sort_values(by=['userId', 'ts'])

# Calculate the time difference between consecutive sessions (session duration)
df['session_diff'] = df.groupby('userId')['ts'].diff().fillna(pd.
    ↳Timedelta(seconds=0))

# Calculate the total session duration (in seconds) for each session
```

```
df['session_duration'] = df['session_diff'].dt.total_seconds()
```

```
[210]: # Dummy Coding

# Page
for click_type in unique_page:
    df[click_type] = np.where(df['page'] == click_type, 1, 0)

# Levels
df['dummy_level'] = np.where(df['level'] == 'paid', 1, 0)
```

4.0.1 Explanation of Features:

- **week:** A new feature to group data by weeks, with each week starting on Monday.
- **session_diff:** Time difference between consecutive sessions, which helps us understand user activity over time.
- **session_duration:** Total session duration in seconds, derived from the session_diff.
- **session_count:** Number of sessions a user has in total.
- **avg_item_in_session:** Average number of items per session.
- **dummy_level:** A binary feature representing whether the user is on a free or paid plan.

```
[10]: display(df.head())
```

```

      status gender  length firstName level lastName  registration
↵userId          ts      auth      page sessionId
↵location itemInSession
↵
↵                                userAgent          song
↵
↵  artist method      week  session_diff session_duration session_count
↵Home NextSong  Thumbs Up  Add to Playlist  Roll Advert  Help  Thumbs Down
↵Upgrade Submit Upgrade  Downgrade  Logout  Add Friend  Settings  About Save
↵Settings Error  Cancel  Cancellation Confirmation  Submit Downgrade  Login
↵Register Submit Registration  dummy_level
403570      200      M      NaN    Camren  free    Walker 2018-07-10 09:30:08
↵1000025 2018-10-02 08:59:29  Logged In      Home    23706 New
↵Haven-Milford, CT      0 "Mozilla/5.0 (Windows NT 6.1; WOW64)
↵AppleWebKit/537.36 (KHTML, like Gecko) Chrome/36.0.1985.143 Safari/537.36"
↵
↵      None      None  GET 2018-10-01 0 days 00:00:00      0.
↵0      122.931172      1      0      0      0      0      0
↵0      0      0      0      0      0      0      0
↵0      0      0      0      0      0      0      0
↵0      0      0      0      0      0      0      0
```

```

404235      200      M  251.97669    Camren free    Walker 2018-07-10 09:30:08
↳1000025 2018-10-02 09:02:44 Logged In NextSong      23706 New
↳Haven-Milford, CT      1 "Mozilla/5.0 (Windows NT 6.1; WOW64)
↳AppleWebKit/537.36 (KHTML, like Gecko) Chrome/36.0.1985.143 Safari/537.36"
↳Underground      1001    PUT 2018-10-01 0 days 00:03:15
↳195.0      122.931172      0      1      0      0      0      0
↳ 0      0      0      0      0      0      0      0
↳ 0      0      0      0      0      0      0      0
↳ 0      0      0      0      0      0      0      0

405134      200      M  223.05914    Camren free    Walker 2018-07-10 09:30:08
↳1000025 2018-10-02 09:06:55 Logged In NextSong      23706 New
↳Haven-Milford, CT      2 "Mozilla/5.0 (Windows NT 6.1; WOW64)
↳AppleWebKit/537.36 (KHTML, like Gecko) Chrome/36.0.1985.143 Safari/537.36"
↳AsesÃÃname Charly GarcÃÃa PUT 2018-10-01 0 days 00:04:11      251.
↳0      122.931172      0      1      0      0      0      0
↳0      0      0      0      0      0      0      0
↳0      0      0      0      0      0      0      0
↳ 0      0      0      0      0      0      0      0

405929      200      M  193.38404    Camren free    Walker 2018-07-10 09:30:08
↳1000025 2018-10-02 09:10:38 Logged In NextSong      23706 New
↳Haven-Milford, CT      3 "Mozilla/5.0 (Windows NT 6.1; WOW64)
↳AppleWebKit/537.36 (KHTML, like Gecko) Chrome/36.0.1985.143 Safari/537.36"
↳Last Nite      The Strokes PUT 2018-10-01 0 days 00:03:43      223.
↳0      122.931172      0      1      0      0      0      0
↳0      0      0      0      0      0      0      0
↳0      0      0      0      0      0      0      0
↳ 0      0      0      0      0      0      0      0

406597      200      M  177.37098    Camren free    Walker 2018-07-10 09:30:08
↳1000025 2018-10-02 09:13:51 Logged In NextSong      23706 New
↳Haven-Milford, CT      4 "Mozilla/5.0 (Windows NT 6.1; WOW64)
↳AppleWebKit/537.36 (KHTML, like Gecko) Chrome/36.0.1985.143 Safari/537.36"
↳ The End      Pearl Jam PUT 2018-10-01 0 days 00:03:13      193.
↳0      122.931172      0      1      0      0      0      0
↳0      0      0      0      0      0      0      0
↳0      0      0      0      0      0      0      0
↳ 0      0      0      0      0      0      0      0

```

5 Defining Churn and Activity

```

[240]: # Identify users who have data for the current week (check for missing rows for
↳each user)
all_users = df['userId'].unique()
weeks_in_df = df['week'].unique()

# Create a DataFrame of all possible combinations of userId and week (even for
↳missing users)

```

```

all_combinations = pd.MultiIndex.from_product([all_users, weeks_in_df],
    ↪names=['userId', 'week'])
df_all_weeks = pd.DataFrame(index=all_combinations).reset_index()

# Create a DataFrame that groups by UserId and Week with the features you are
    ↪interested in
df_actuals = df.groupby(['userId', 'week']).agg(
    session_count = ('itemInSession', 'mean'),
    session_duration = ('session_duration', 'sum'),
    dummy_level = ('dummy_level', 'first')
)
# Merge with the original data to identify users missing for a given week
df_all_weeks = df_all_weeks.merge(df_actuals, on=['userId', 'week'], how='left')

# Sort by userId and week to handle previous week properly
df_all_weeks = df_all_weeks.sort_values(by=['userId', 'week'])

# Create "prev_week_activity" to track if the user was active in the previous
    ↪week
df_all_weeks['prev_week_activity'] = df_all_weeks.groupby('userId',
    ↪group_keys=False)['session_count'].shift(1).fillna(0)
df_all_weeks['prev_week_session_duration'] = df_all_weeks.groupby('userId',
    ↪group_keys=False)['session_duration'].shift(1).fillna(0)
df_all_weeks['prev_dummy_level'] = df_all_weeks.groupby('userId',
    ↪group_keys=False)['dummy_level'].shift(1).fillna(method='ffill').fillna(0)

# Fill NaN session_count, prev_week_activity, and session_duration with 0
df_all_weeks['session_count'] = df_all_weeks['session_count'].fillna(0)
df_all_weeks['session_duration'] = df_all_weeks['session_duration'].fillna(0)

# Fill dummy_level NaN with forward-fill (ffill)
df_all_weeks['dummy_level'] = df_all_weeks['dummy_level'].fillna(method='ffill')

# Define churn: users who were active in the previous week but are inactive now
df_all_weeks['churn'] = np.where(
    ((df_all_weeks['session_count'] == 0) & (df_all_weeks['prev_week_activity']
    ↪> 0)),
    1,
    0
)

```

/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/1790987056.py:25: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```

df_all_weeks['prev_dummy_level'] = df_all_weeks.groupby('userId',
group_keys=False)['dummy_level'].shift(1).fillna(method='ffill').fillna(0)

```

```
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/1790987056.py:3
2: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a
future version. Use obj.ffill() or obj.bfill() instead.
```

```
df_all_weeks['dummy_level'] =
df_all_weeks['dummy_level'].fillna(method='ffill')
```

```
[ ]: # Sort for right calculations
df = df.sort_values(by=['ts'], ascending=True)

# Group by Days
df['day'] = df['ts'].dt.date

df_all_days = df.groupby(['userId', 'week', 'day']).agg(
    session_count=('itemInSession', 'mean'),
    session_duration=('session_duration', 'sum'),
    dummy_level=('dummy_level', 'first')
)

# Ensure daily data is sorted
df_all_days = df_all_days.sort_values(by=['userId', 'day'])

# Define rolling window sizes in days
window_sizes = [7, 14, 21, 28] # Rolling windows of 1, 2, 3, and 4 weeks

# Apply rolling difference calculations
for window in window_sizes:
    df_all_days[f'rolling_diff_{window}_days'] = (
        df_all_days.groupby('userId')['session_count']
        .rolling(window=window, min_periods=1)
        .sum()
        .diff()
        .reset_index(level=0, drop=True)
    )

# Aggregate daily rolling differences to weekly level (take last value in each
↪ week)
df_weekly_rolling_diffs = df_all_days.groupby(['userId', 'week'])[[
    f'rolling_diff_{window}_days' for window in window_sizes
]].last().reset_index()

# Merge corrected rolling diffs back into df_all_weeks
df_all_weeks = df_all_weeks.merge(df_weekly_rolling_diffs, on=['userId', ↪
↪ 'week'], how='left')
```

```

# Sort before shifting
df_all_weeks = df_all_weeks.sort_values(by=['userId', 'week'])

# Apply shift for previous week rolling diffs
for window in window_sizes:
    df_all_weeks[f'prev_week_rolling_diff_{window}_days'] = (
        df_all_weeks.groupby('userId',
                               group_keys=False)[f'rolling_diff_{window}_days']
        .shift(1) # Shift values down by one week per user
        .fillna(0)
    )

# # Display final dataset

print(df_all_weeks)

```

	userId	week	session_count	session_duration	dummy_level	prev_week_activity	prev_week_session_duration	prev_dummy_level	churn	
			rolling_diff_7_days	rolling_diff_14_days	rolling_diff_21_days	rolling_diff_28_days	prev_week_rolling_diff_7_days	prev_week_rolling_diff_14_days	prev_week_rolling_diff_21_days	prev_week_rolling_diff_28_days
0	1000025	2018-10-01	171.821705	482224.0	0.0	0.000000	0.0	0.0	0	68.703704
			0.0	0.0	0	68.703704	68.703704	68.703704	68.703704	0.000000
			0.000000	0.000000	0.000000	0.000000				0.000000
1	1000025	2018-10-08	66.317416	350667.0	1.0	171.821705	482224.0	0.0	0	68.703704
			482224.0	0.0	0	-27.569620	94.000000	94.000000	94.000000	68.703704
			94.000000	94.000000	94.000000	68.703704	68.703704	68.703704	68.703704	68.703704
2	1000025	2018-10-15	90.768097	591125.0	1.0	66.317416	350667.0	1.0	0	120.773684
			350667.0	1.0	0	120.773684	130.273684	130.273684	130.273684	130.273684
			130.273684	130.273684	130.273684	-27.569620	94.000000	94.000000	94.000000	94.000000
3	1000025	2018-10-22	0.000000	0.0	1.0	90.768097	591125.0	1.0	1	NaN
			591125.0	1.0	1	NaN	NaN	NaN	NaN	NaN
			NaN	NaN	NaN	120.773684	130.273684	130.273684	130.273684	130.273684
4	1000025	2018-10-29	0.000000	0.0	1.0	0.000000	0.0	1.0	0	NaN
			0.0	1.0	0	NaN	NaN	NaN	NaN	NaN
			0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
...	...	...	...	...	...	...	...	...	...	...

```

...
...
...
200497 1999996 2018-10-29 15.950820 199341.0 0.0
39.618799 1143249.0 0.0 0
-50.707407 17.500000 17.500000 17.500000
20.461538 24.461538 24.461538
24.461538
200498 1999996 2018-11-05 25.627451 726498.0 0.0
15.950820 199341.0 0.0 0
21.627451 25.627451 25.627451 25.627451
-50.707407 17.500000 17.500000
17.500000
200499 1999996 2018-11-12 3.500000 455443.0 0.0
25.627451 726498.0 0.0 0
-25.836207 -6.000000 3.500000 3.500000
21.627451 25.627451 25.627451
25.627451
200500 1999996 2018-11-19 0.000000 0.0 0.0
3.500000 455443.0 0.0 1
NaN NaN NaN NaN
-25.836207 -6.000000 3.500000
3.500000
200501 1999996 2018-11-26 6.000000 1225167.0 0.0
0.000000 0.0 0.0 0
-12.555556 5.500000 6.000000 6.000000
0.000000 0.000000 0.000000
0.000000

```

[200502 rows x 17 columns]

```
[192]: display(df_all_weeks)
display(df_actuals)
```

```

      userId      week  session_count  session_duration  dummy_level  ↵
↵prev_week_activity  prev_week_session_duration  prev_dummy_level  churn  ↵
↵rolling_diff_1_weeks  rolling_diff_2_weeks  rolling_diff_3_weeks  ↵
↵rolling_diff_4_weeks  rolling_diff_5_weeks  rolling_diff_6_weeks  ↵
↵rolling_diff_7_weeks  rolling_diff_8_weeks
127487 1638340 2018-10-15      0.000000      0.0      0.0      ↵
↵      8.000000      682077.0      0.0      1      ↵
↵7.500000      8.000000      8.000000      8.000000      8.000000      ↵
↵      8.000000      8.000000      8.000000      8.000000      8.
↵000000

```

123483	1617370	2018-10-22	0.000000	0.0	0.0	□
↪	36.746269		27440.0	0.0	1	□
↪	36.746269	36.746269		36.746269	36.746269	□
↪	36.746269	36.746269		36.746269	36.	
↪	746269					
179224	1892578	2018-11-19	0.000000	0.0	0.0	□
↪	27.710145		586791.0	0.0	1	□
↪	-4.833973	27.710145		24.210145	0.539690	□
↪	27.710145	27.710145		27.710145	27.	
↪	710145					
72079	1358034	2018-11-19	0.000000	0.0	1.0	□
↪	70.799270		650214.0	1.0	1	□
↪	19.985417	-86.630217		43.864964	61.694007	□
↪	16.018782	47.799270		70.799270	70.	
↪	799270					
49255	1239501	2018-11-19	0.000000	0.0	0.0	□
↪	11.000000		370161.0	0.0	1	□
↪	10.000000	11.000000		11.000000	11.000000	□
↪	11.000000	11.000000		11.000000	11.	
↪	000000					
...	...	...	...	...	...	□
↪		...	...	...	...	□
↪	...	...	...	...	...	□
↪	...	...	...	...	...	
72040	1357698	2018-10-29	59.946121	617104.0	1.0	□
↪	72.725333		592484.0	1.0	0	□
↪	-21.913300	66.689619		55.581103	72.725333	□
↪	72.725333	72.725333		72.725333	72.	
↪	725333					
72041	1357698	2018-11-05	164.284264	604744.0	1.0	□
↪	59.946121		617104.0	1.0	0	□
↪	-12.779213	-34.692513		53.910406	42.801890	□
↪	59.946121	59.946121		59.946121	59.	
↪	946121					
72042	1357698	2018-11-12	99.775126	586588.0	0.0	□
↪	164.284264		604744.0	1.0	0	□
↪	104.338143	91.558931		69.645631	158.248550	□
↪	147.140033	164.284264		164.284264	164.	
↪	284264					
72043	1357698	2018-11-19	195.478462	622826.0	1.0	□
↪	99.775126		586588.0	0.0	0	□
↪	-64.509138	39.829005		27.049792	5.136492	□
↪	93.739411	82.630895		99.775126	99.	
↪	775126					

200501	1999996	2018-11-26	6.000000	1225167.0	0.0	
↪	0.000000		0.0	0.0	0	
↪	-3.500000	-25.627451		-15.950820	-39.618799	
↪		-4.000000	-3.500000	-8.600000		0.
↪	000000					

[200502 rows x 17 columns]

		session_count	session_duration	dummy_level
userId	week			
1000025	2018-10-01	171.821705	482224.0	0
	2018-10-08	66.317416	350667.0	1
	2018-10-15	90.768097	591125.0	1
1000035	2018-10-01	4.500000	58000.0	0
	2018-10-08	34.097561	459822.0	0
...		...	...	...
1999996	2018-10-22	39.618799	1143249.0	0
	2018-10-29	15.950820	199341.0	0
	2018-11-05	25.627451	726498.0	0
	2018-11-12	3.500000	455443.0	0
	2018-11-26	6.000000	1225167.0	0

[121443 rows x 3 columns]

```
[242]: # Additional Feature Engineering

# Ensure data is sorted by userId and week before applying rolling operations
df_all_weeks = df_all_weeks.sort_values(by=['userId', 'week'])

# Define rolling window sizes
window_sizes = [1, 2, 3, 4, 5, 6, 7, 8]

# Apply rolling difference calculations for multiple window sizes
for window in window_sizes:
    df_all_weeks[f'rolling_diff_{window}_weeks'] = (
        df_all_weeks.groupby('userId')['prev_week_activity']
            .rolling(window=window, min_periods=1)
            .sum()
            .diff()
            .reset_index(level=0, drop=True) # Ensures proper alignment after
↪groupby
    )
    # Replace NaN values with 0
    df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
```

/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:19: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:1
9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:1
9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:1
9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:1
9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:1
9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:1
9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
/var/folders/fl/6yh7c335749bvjp9pb0rj8fw0000gn/T/ipykernel_20340/4092923539.py:1
9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using

'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df_all_weeks[f'rolling_diff_{window}_weeks'].fillna(0, inplace=True)
```

```
[243]: df_all_weeks = df_all_weeks.sort_values(by='churn', ascending=False)
display(df_all_weeks.head(100))
```

```
print(df_all_weeks['churn'].dtypes)
print(df_all_weeks['churn'].unique())
```

```
# Check for NaN.
```

```
for column in df_all_weeks.columns:
    nan_count = df_all_weeks[column].isna().sum()
    print(f"Column '{column}' has {nan_count} NaN values.")
```

```

      userId      week  session_count  session_duration  dummy_level  ↵
↵prev_week_activity  prev_week_session_duration  prev_dummy_level  churn  ↵
↵rolling_diff_7_days  rolling_diff_14_days  rolling_diff_21_days  ↵
↵rolling_diff_28_days  prev_week_rolling_diff_7_days  ↵
↵prev_week_rolling_diff_14_days  prev_week_rolling_diff_21_days  ↵
↵prev_week_rolling_diff_28_days  rolling_diff_1_weeks  rolling_diff_2_weeks  ↵
↵rolling_diff_3_weeks  rolling_diff_4_weeks  rolling_diff_5_weeks  ↵
↵rolling_diff_6_weeks  rolling_diff_7_weeks  rolling_diff_8_weeks
127487  1638340  2018-10-15          0.0          0.0          0.0      ↵
↵  8.000000          682077.0          0.0          1      ↵
↵  NaN          NaN          NaN          NaN      ↵
↵          8.000000          8.000000      ↵
↵  8.000000          8.000000          7.500000      8.
↵000000          8.000000          8.000000          8.000000      ↵
↵          8.000000          8.000000          8.000000
123483  1617370  2018-10-22          0.0          0.0          0.0      ↵
↵  36.746269          27440.0          0.0          1      ↵
↵  NaN          NaN          NaN          NaN      ↵
↵          67.000000          67.000000      ↵
↵67.000000          67.000000          36.746269      36.
↵746269          36.746269          36.746269          36.746269      ↵
↵          36.746269          36.746269          36.746269
179224  1892578  2018-11-19          0.0          0.0          0.0      ↵
↵  27.710145          586791.0          0.0          1      ↵
↵  NaN          NaN          NaN          NaN      ↵
↵          53.500000          53.500000      ↵
↵53.500000          53.500000          -4.833973      27.
↵710145          24.210145          0.539690          27.710145      ↵
↵          27.710145          27.710145          27.710145

```

72079	1358034	2018-11-19	0.0	0.0	1.0	
↪	70.799270		650214.0	1.0	1	
↪	NaN		NaN	NaN	NaN	
↪		24.889189		94.389189		
↪	94.389189		94.389189	19.985417		-86.
↪	630217		43.864964	61.694007	16.018782	
↪	47.799270		70.799270	70.799270		
49255	1239501	2018-11-19	0.0	0.0	0.0	
↪	11.000000		370161.0	0.0	1	
↪	NaN		NaN	NaN	NaN	
↪		11.000000		11.000000		
↪	11.000000		11.000000	10.000000		11.
↪	000000		11.000000	11.000000	11.000000	
↪	11.000000		11.000000	11.000000		
...	...	...	...	...	...	...
↪		...	...	...	...	
↪	...	...	...	...	...	
↪		...	...	...	...	
↪	...	...	...	...	...	
↪	...	...	...	...	...	
↪	...	...	...	...	...	
179338	1893052	2018-10-29	0.0	0.0	0.0	
↪	62.479730		767589.0	0.0	1	
↪	NaN		NaN	NaN	NaN	
↪		9.000000		9.000000		
↪	9.000000		9.000000	62.479730		56.
↪	979730		40.592633	62.479730	62.479730	
↪	62.479730		62.479730	62.479730		
24706	1122985	2018-10-08	0.0	0.0	1.0	
↪	55.347368		57989.0	1.0	1	
↪	NaN		NaN	NaN	NaN	
↪		98.014286		98.014286		
↪	98.014286		98.014286	55.347368		55.
↪	347368		55.347368	55.347368	55.347368	
↪	55.347368		55.347368	55.347368		
89956	1448255	2018-10-08	0.0	0.0	0.0	
↪	8.000000		4013.0	0.0	1	
↪	NaN		NaN	NaN	NaN	
↪		-280.536042		-280.536042		
↪	-280.536042		-280.536042	8.000000		
↪	8.000000		8.000000	8.000000	8.000000	
↪	8.000000		8.000000	8.000000		

24746	1123156	2018-11-05	0.0	0.0	1.0	
↪ 68.583799			687512.0	1.0	1	
↪ NaN		NaN		NaN		NaN
↪	22.000000			-9.828205		
↪ 44.500000		44.500000		22.509541		31.
↪ 188276	-5.431276		-21.982509		68.583799	
↪ 68.583799		68.583799		68.583799		
123438	1617154	2018-10-22	0.0	0.0	1.0	
↪ 7.000000			69378.0	1.0	1	
↪ NaN		NaN		NaN		NaN
↪	7.000000			7.000000		
↪ 7.000000		7.000000		-53.488688		-0.
↪ 500000	7.000000		7.000000		7.000000	
↪ 7.000000		7.000000		7.000000		

[100 rows x 25 columns]

int64

[1 0]

Column 'userId' has 0 NaN values.

Column 'week' has 0 NaN values.

Column 'session_count' has 0 NaN values.

Column 'session_duration' has 0 NaN values.

Column 'dummy_level' has 0 NaN values.

Column 'prev_week_activity' has 0 NaN values.

Column 'prev_week_session_duration' has 0 NaN values.

Column 'prev_dummy_level' has 0 NaN values.

Column 'churn' has 0 NaN values.

Column 'rolling_diff_7_days' has 79059 NaN values.

Column 'rolling_diff_14_days' has 79059 NaN values.

Column 'rolling_diff_21_days' has 79059 NaN values.

Column 'rolling_diff_28_days' has 79059 NaN values.

Column 'prev_week_rolling_diff_7_days' has 0 NaN values.

Column 'prev_week_rolling_diff_14_days' has 0 NaN values.

Column 'prev_week_rolling_diff_21_days' has 0 NaN values.

Column 'prev_week_rolling_diff_28_days' has 0 NaN values.

Column 'rolling_diff_1_weeks' has 0 NaN values.

Column 'rolling_diff_2_weeks' has 0 NaN values.

Column 'rolling_diff_3_weeks' has 0 NaN values.

Column 'rolling_diff_4_weeks' has 0 NaN values.

Column 'rolling_diff_5_weeks' has 0 NaN values.

Column 'rolling_diff_6_weeks' has 0 NaN values.

Column 'rolling_diff_7_weeks' has 0 NaN values.

Column 'rolling_diff_8_weeks' has 0 NaN values.

```
[148]: display(df_all_weeks[df_all_weeks['userId'] == 1358034].
↪sort_values('week').head(100))
```

	userId	week	session_count	session_duration	dummy_level	␣
↪	prev_week_activity	prev_week_session_duration	prev_dummy_level	churn	␣	
↪	rolling_diff_3_weeks	rolling_diff_5_weeks	rolling_diff_10_weeks	␣		
↪	rolling_diff_15_weeks	rolling_diff_20_weeks	rolling_diff_1_weeks	␣		
↪	rolling_diff_2_weeks	rolling_diff_4_weeks	rolling_diff_6_weeks	␣		
↪	rolling_diff_7_weeks	rolling_diff_8_weeks				
72072	1358034	2018-10-01	23.000000	10847.0	0.0	␣
↪	0.000000		0.0	1.0	0	-188.
↪	940993	-188.940993		-231.940993	-231.940993	␣
↪		-231.940993	-79.438710	-188.940993		-188.
↪	940993	-209.940993	-209.940993	-231.940993		
72073	1358034	2018-10-08	54.780488	358279.0	0.0	␣
↪	23.000000		10847.0	0.0	0	23.
↪	000000	23.000000		23.000000	23.000000	␣
↪		23.000000	23.000000	23.000000		23.
↪	000000	23.000000		23.000000	23.000000	
72074	1358034	2018-10-15	9.105263	522903.0	1.0	␣
↪	54.780488		358279.0	0.0	0	54.
↪	780488	54.780488		54.780488	54.780488	␣
↪		54.780488	31.780488	54.780488		54.
↪	780488	54.780488		54.780488	54.780488	
72075	1358034	2018-10-22	26.934307	1004345.0	1.0	␣
↪	9.105263		522903.0	1.0	0	9.
↪	105263	9.105263		9.105263	9.105263	␣
↪		9.105263	-45.675225	-13.894737		9.
↪	105263	9.105263		9.105263	9.105263	
72076	1358034	2018-10-29	157.429487	433527.0	1.0	␣
↪	26.934307		1004345.0	1.0	0	3.
↪	934307	26.934307		26.934307	26.934307	␣
↪		26.934307	17.829043	-27.846181		26.
↪	934307	26.934307		26.934307	26.934307	
72077	1358034	2018-11-05	50.813853	686717.0	1.0	␣
↪	157.429487		433527.0	1.0	0	102.
↪	648999	157.429487		157.429487	157.429487	␣
↪		157.429487	130.495181	148.324224		134.
↪	429487	157.429487		157.429487	157.429487	
72078	1358034	2018-11-12	70.799270	650214.0	1.0	␣
↪	50.813853		686717.0	1.0	0	41.
↪	708590	27.813853		50.813853	50.813853	␣
↪		50.813853	-106.615634	23.879546		-3.
↪	966635	50.813853		50.813853	50.813853	
72079	1358034	2018-11-19	0.000000	0.0	1.0	␣
↪	70.799270		650214.0	1.0	1	43.
↪	864964	16.018782		70.799270	70.799270	␣
↪		70.799270	19.985417	-86.630217		61.
↪	694007	47.799270		70.799270	70.799270	

```

72080  1358034  2018-11-26      56.880851      1194396.0      1.0      0.000000      0.0      1.0      0      -157.
↳ 429487      -9.105263      0.000000      -50.813853      0.000000      -26.
↳ 934307      -54.780488      -23.000000      0.000000

```

```
[53]: # Testing purposes:
```

```

df_first_weeks = df_all_weeks[(df_all_weeks['week'] == '2018-10-08') |
↳ (df_all_weeks['week'] == '2018-10-15') | (df_all_weeks['week'] ==
↳ '2018-10-22') | (df_all_weeks['week'] == '2018-10-29')]

```

```
[149]: print(df_first_weeks['churn'].value_counts(normalize=True))
```

```

churn
0    0.998815
1    0.001185
Name: proportion, dtype: float64

```

6 Training the Churn Model (Rolling Prediction)

```

[244]: def rolling_churn_prediction(df, features, label, model=None):
        """
        Perform week-to-week churn predictions using a rolling window approach.

        Parameters:
        - df: Aggregated user activity dataframe with churn labels.
        - features: List of feature column names.
        - label: Target column name (churn).
        - model: Machine learning model (default: Random Forest with
↳ class_weight='balanced')

        Returns:
        - None (prints classification reports for each week).
        """
        if model is None:
            model = RandomForestClassifier(n_estimators=100, random_state=42,
↳ class_weight='balanced') # Handle class imbalance

        print(f"Model type: {type(model)}")

        weeks = sorted(df['week'].unique()) # Get sorted list of weeks
        f1_scores = []
        week_labels = []

        for i in range(len(weeks) - 1): # Stop at second-to-last week (since we
↳ predict next week)

```



```

train_weeks = weeks[i:i + 1] # Train up to current week
test_week = weeks[i + 1] # Predict for the next week

train = df[df['week'].isin(train_weeks)]
test = df[df['week'] == test_week]

if test.empty: # Skip if no data for the next week
    continue

# Convert categorical features to numeric
X_train = train[features]
y_train = train[label]
X_test = test[features]
y_test = test[label]

# X_train, X_test = X_train.align(X_test, join="left", axis=1,
↪fill_value=0)

# Train the model
model.fit(X_train, y_train)

# Predict churn
y_pred = model.predict(X_test)

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

# Compute specificity
specificity = tn / (tn + fp) if (tn + fp) != 0 else 0

# Generate classification report as a dictionary
report_dict = classification_report(y_test, y_pred, output_dict=True,
↪zero_division=0)

# Add specificity to the report
report_dict['0']['specificity'] = specificity

# Convert report to DataFrame for better display
report_df = pd.DataFrame(report_dict).transpose()

# Print evaluation results for each week
print(f"Week {test_week}:")
print(report_df)

# Plot confusion matrix

```

```

    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['No_
↳Churn', 'Churn'])
    disp.plot(cmap='Blues')
    plt.title(f"Confusion Matrix - Week {test_week}")
    plt.show()

    # F1 Score tracking
    f1 = f1_score(y_test, y_pred, zero_division=0)
    f1_scores.append(f1)
    week_labels.append(str(test_week))

    print("-" * 50)

    # Plot F1-scores over weeks
    if f1_scores:
        plt.figure(figsize=(10, 6))
        plt.plot(week_labels, f1_scores, marker='o')
        plt.title('F1-Score for Churn Prediction Over Weeks')
        plt.xlabel('Week')
        plt.ylabel('F1-Score')
        plt.xticks(rotation=45)
        plt.grid(True)
        plt.tight_layout()
        plt.show()
    else:
        print("No evaluations performed (possibly no data after filtering).")

```

```

[ ]: model = lgb.LGBMClassifier(random_state=42)
features = ['prev_week_activity',
            'prev_week_session_duration',
            'prev_dummy_level',
            'rolling_diff_1_weeks',
            'rolling_diff_2_weeks',
            'rolling_diff_3_weeks',
            'rolling_diff_4_weeks',
            'rolling_diff_5_weeks',
            'rolling_diff_6_weeks',
            'rolling_diff_7_weeks',
            'rolling_diff_8_weeks',
            'prev_week_rolling_diff_7_days',
            'prev_week_rolling_diff_14_days',
            'prev_week_rolling_diff_21_days',
            'prev_week_rolling_diff_28_days'
            ]
label = 'churn'


```


the split requirements

the split requirements


the split requirements

Week 2018-10-08 00:00:00:




```
-----
[LightGBM] [Info] Number of positive: 3961, number of negative: 18317
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.001477 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 3572
[LightGBM] [Info] Number of data points in the train set: 22278, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.177799 -> initscore=-1.531333
[LightGBM] [Info] Start training from score -1.531333
Week 2018-10-15 00:00:00:
```

	precision	recall	f1-score	support	specificity
0	0.833175	0.994316	0.906641	18474.000000	0.994316
1	0.545455	0.033123	0.062454	3804.000000	NaN
accuracy	0.830191	0.830191	0.830191	0.830191	0.830191
macro avg	0.689315	0.513720	0.484547	22278.000000	NaN
weighted avg	0.784046	0.830191	0.762495	22278.000000	NaN




```
-----
[LightGBM] [Info] Number of positive: 3804, number of negative: 18474
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.001069 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 3572
[LightGBM] [Info] Number of data points in the train set: 22278, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.170751 -> initscore=-1.580311
[LightGBM] [Info] Start training from score -1.580311
Week 2018-10-22 00:00:00:
```

	precision	recall	f1-score	support	specificity
0	0.845473	0.982956	0.909046	18658.000000	0.982956
1	0.457338	0.074033	0.127437	3620.000000	NaN
accuracy	0.835263	0.835263	0.835263	0.835263	0.835263
macro avg	0.651405	0.528495	0.518241	22278.000000	NaN
weighted avg	0.782404	0.835263	0.782041	22278.000000	NaN




```
-----
[LightGBM] [Info] Number of positive: 3620, number of negative: 18658
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000529 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3572
[LightGBM] [Info] Number of data points in the train set: 22278, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.162492 -> initscore=-1.639801
[LightGBM] [Info] Start training from score -1.639801
Week 2018-10-29 00:00:00:
```

	precision	recall	f1-score	support	specificity
0	0.851191	0.969907	0.906679	18642.000000	0.969907
1	0.458494	0.130638	0.203339	3636.000000	NaN
accuracy	0.832929	0.832929	0.832929	0.832929	0.832929
macro avg	0.654843	0.550272	0.555009	22278.000000	NaN
weighted avg	0.787099	0.832929	0.791887	22278.000000	NaN




```
-----
[LightGBM] [Info] Number of positive: 3636, number of negative: 18642
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000453 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3572
[LightGBM] [Info] Number of data points in the train set: 22278, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.163210 -> initscore=-1.634533
[LightGBM] [Info] Start training from score -1.634533
Week 2018-11-05 00:00:00:
```

	precision	recall	f1-score	support	specificity
0	0.865736	0.956885	0.909031	18787.000000	0.956885
1	0.464640	0.201375	0.280975	3491.000000	NaN
accuracy	0.838495	0.838495	0.838495	0.838495	0.838495
macro avg	0.665188	0.579130	0.595003	22278.000000	NaN
weighted avg	0.802883	0.838495	0.810614	22278.000000	NaN




```
-----
[LightGBM] [Info] Number of positive: 3491, number of negative: 18787
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000370 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3572
[LightGBM] [Info] Number of data points in the train set: 22278, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.156702 -> initscore=-1.682977
[LightGBM] [Info] Start training from score -1.682977
Week 2018-11-12 00:00:00:
```

	precision	recall	f1-score	support	specificity
0	0.871828	0.960176	0.913872	18858.0000	0.960176
1	0.502319	0.221637	0.307567	3420.0000	NaN
accuracy	0.846800	0.846800	0.846800	0.8468	0.846800
macro avg	0.687074	0.590907	0.610720	22278.0000	NaN
weighted avg	0.815103	0.846800	0.820795	22278.0000	NaN





```
-----
[LightGBM] [Info] Number of positive: 3420, number of negative: 18858
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000288 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3572
[LightGBM] [Info] Number of data points in the train set: 22278, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.153515 -> initscore=-1.707297
[LightGBM] [Info] Start training from score -1.707297
Week 2018-11-19 00:00:00:
```

	precision	recall	f1-score	support	specificity
0	0.859743	0.968185	0.910747	18639.000000	0.968185
1	0.539596	0.190987	0.282119	3639.000000	NaN
accuracy	0.841234	0.841234	0.841234	0.841234	0.841234
macro avg	0.699670	0.579586	0.596433	22278.000000	NaN
weighted avg	0.807448	0.841234	0.808064	22278.000000	NaN



```
-----
[LightGBM] [Info] Number of positive: 3639, number of negative: 18639
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000406 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3572
[LightGBM] [Info] Number of data points in the train set: 22278, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.163345 -> initscore=-1.633547
[LightGBM] [Info] Start training from score -1.633547
Week 2018-11-26 00:00:00:
```

	precision	recall	f1-score	support	specificity
0	0.872461	0.947239	0.908313	18726.000000	0.947239
1	0.492553	0.269989	0.348791	3552.000000	NaN
accuracy	0.839258	0.839258	0.839258	0.839258	0.839258
macro avg	0.682507	0.608614	0.628552	22278.000000	NaN
weighted avg	0.811888	0.839258	0.819103	22278.000000	NaN



```

-----
KeyError                                Traceback (most recent call last)
Cell In[245], line 25
    22 # Apply shift for previous week rolling diffs
    23 for window in window_sizes:
    24     df_all_weeks[f'prev_week_rolling_diff_{window}_days'] = (
---> 25         □
    26     df_all_weeks.groupby('userId', group_keys=False)[f'rolling_diff_{window}_days']
    27         .shift(1) # Shift values down by one week per user
    28         .fillna(0)
    29     )

File ~/Documents/Seminar/.venv/lib/python3.13/site-packages/pandas/core/groupby/
generic.py:1951, in DataFrameGroupBy.__getitem__(self, key)
    1944 if isinstance(key, tuple) and len(key) > 1:
    1945     # if len == 1, then it becomes a SeriesGroupBy and this is actually
    1946     # valid syntax, so don't raise
    1947     raise ValueError(
    1948         "Cannot subset columns with a tuple with more than one element. "
    1949         "Use a list instead."
    1950     )
-> 1951 return super().__getitem__(key)

File ~/Documents/Seminar/.venv/lib/python3.13/site-packages/pandas/core/base.py
244, in SelectionMixin.__getitem__(self, key)
    242 else:
    243     if key not in self.obj:
--> 244         raise KeyError(f"Column not found: {key}")
    245     ndim = self.obj[key].ndim
    246     return self._getitem(key, ndim=ndim)

KeyError: 'Column not found: rolling_diff_1_days'

```

```

[246]: # Fitting a XGB model
model = XGBClassifier(n_estimators=100, random_state=42, base_score=0.5, □
    scale_pos_weight=10)
features = ['prev_week_activity',
            'prev_week_session_duration',
            'prev_dummy_level',
            'rolling_diff_1_weeks',
            'rolling_diff_2_weeks',
            'rolling_diff_3_weeks',
            'rolling_diff_4_weeks',
            'rolling_diff_5_weeks',

```



```

        'rolling_diff_6_weeks',
        'rolling_diff_7_weeks',
        'rolling_diff_8_weeks',
        'prev_week_rolling_diff_7_days',
        'prev_week_rolling_diff_14_days',
        'prev_week_rolling_diff_21_days',
        'prev_week_rolling_diff_28_days'
    ]
    label = 'churn'


    rolling_churn_prediction(df = df_all_weeks, features=features, label=label,
                             ↪model=model)

```

Model type: <class 'xgboost.sklearn.XGBClassifier'>


Week 2018-10-08 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.822201	1.000000	0.902426	18317.000000	1.000000
1	0.000000	0.000000	0.000000	3961.000000	NaN
accuracy	0.822201	0.822201	0.822201	0.822201	0.822201
macro avg	0.411101	0.500000	0.451213	22278.000000	NaN
weighted avg	0.676015	0.822201	0.741976	22278.000000	NaN




Week 2018-10-15 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.886814	0.843131	0.864421	18474.00000	0.843131
1	0.385235	0.477392	0.426391	3804.00000	NaN
accuracy	0.780680	0.780680	0.780680	0.78068	0.780680
macro avg	0.636025	0.660262	0.645406	22278.00000	NaN
weighted avg	0.801169	0.780680	0.789627	22278.00000	NaN




Week 2018-10-22 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.917835	0.716047	0.804480	18658.000000	0.716047
1	0.313908	0.669613	0.427438	3620.000000	NaN
accuracy	0.708502	0.708502	0.708502	0.708502	0.708502
macro avg	0.615871	0.692830	0.615959	22278.000000	NaN
weighted avg	0.819701	0.708502	0.743214	22278.000000	NaN




Week 2018-10-29 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.932401	0.740639	0.825531	18642.000000	0.740639
1	0.352744	0.724697	0.474518	3636.000000	NaN
accuracy	0.738038	0.738038	0.738038	0.738038	0.738038
macro avg	0.642573	0.732668	0.650024	22278.000000	NaN
weighted avg	0.837795	0.738038	0.768242	22278.000000	NaN




Week 2018-11-05 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.931293	0.766221	0.840731	18787.000000	0.766221
1	0.356106	0.695789	0.471102	3491.000000	NaN
accuracy	0.755184	0.755184	0.755184	0.755184	0.755184
macro avg	0.643700	0.731005	0.655916	22278.000000	NaN
weighted avg	0.841160	0.755184	0.782810	22278.000000	NaN




Week 2018-11-12 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.940776	0.761481	0.841686	18858.000000	0.761481
1	0.358711	0.735673	0.482270	3420.000000	NaN
accuracy	0.757519	0.757519	0.757519	0.757519	0.757519
macro avg	0.649743	0.748577	0.661978	22278.000000	NaN
weighted avg	0.851420	0.757519	0.786510	22278.000000	NaN





Week 2018-11-19 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.930923	0.784484	0.851453	18639.000000	0.784484
1	0.388678	0.701841	0.500294	3639.000000	NaN
accuracy	0.770985	0.770985	0.770985	0.770985	0.770985
macro avg	0.659800	0.743163	0.675873	22278.000000	NaN
weighted avg	0.842350	0.770985	0.794093	22278.000000	NaN



Week 2018-11-26 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.943244	0.785432	0.857135	18726.000000	0.785432
1	0.398953	0.750845	0.521051	3552.000000	NaN
accuracy	0.779917	0.779917	0.779917	0.779917	0.779917
macro avg	0.671098	0.768138	0.689093	22278.000000	NaN
weighted avg	0.856462	0.779917	0.803549	22278.000000	NaN





```
[247]: print(df_all_weeks[features + [label]].corr()[label].
        ↪sort_values(ascending=False))
```

```
churn                                1.000000
prev_week_session_duration           0.273998
rolling_diff_4_weeks                 0.156740
rolling_diff_5_weeks                 0.155946
rolling_diff_3_weeks                 0.155398
rolling_diff_2_weeks                 0.153361
rolling_diff_6_weeks                 0.152086
rolling_diff_7_weeks                 0.149336
rolling_diff_8_weeks                 0.147695
prev_week_activity                   0.142876
rolling_diff_1_weeks                 0.138172
prev_dummy_level                     -0.037114
prev_week_rolling_diff_7_days        -0.040808
prev_week_rolling_diff_14_days       -0.069184
prev_week_rolling_diff_21_days       -0.079749
prev_week_rolling_diff_28_days       -0.082911
Name: churn, dtype: float64
```

```
[248]: # Fitting a Random forest model
features = ['prev_week_activity',
            'prev_week_session_duration',
            'prev_dummy_level',
            'rolling_diff_1_weeks',
            'rolling_diff_2_weeks',
            'rolling_diff_3_weeks',
            'rolling_diff_4_weeks',
            'rolling_diff_5_weeks',
            'rolling_diff_6_weeks',
            'rolling_diff_7_weeks',
            'rolling_diff_8_weeks',
            'prev_week_rolling_diff_7_days',
            'prev_week_rolling_diff_14_days',
            'prev_week_rolling_diff_21_days',
            'prev_week_rolling_diff_28_days'
          ]
label = 'churn'
rolling_churn_prediction(df = df_all_weeks, features=features, label=label,
        ↪model=None)
```


```
Model type: <class 'sklearn.ensemble._forest.RandomForestClassifier'>
Week 2018-10-08 00:00:00:
      precision    recall  f1-score   support  specificity
0          0.822201    1.000000    0.902426   18317.000000        1.000000
```

1	0.000000	0.000000	0.000000	3961.000000	NaN
accuracy	0.822201	0.822201	0.822201	0.822201	0.822201
macro avg	0.411101	0.500000	0.451213	22278.000000	NaN
weighted avg	0.676015	0.822201	0.741976	22278.000000	NaN




Week 2018-10-15 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.837437	0.984898	0.905201	18474.000000	0.984898
1	0.493648	0.071504	0.124914	3804.000000	NaN
accuracy	0.828934	0.828934	0.828934	0.828934	0.828934
macro avg	0.665543	0.528201	0.515058	22278.000000	NaN
weighted avg	0.778735	0.828934	0.771966	22278.000000	NaN




Week 2018-10-22 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.847240	0.981241	0.909330	18658.000000	0.981241
1	0.476831	0.088122	0.148753	3620.000000	NaN
accuracy	0.836116	0.836116	0.836116	0.836116	0.836116
macro avg	0.662035	0.534681	0.529041	22278.000000	NaN
weighted avg	0.787051	0.836116	0.785742	22278.000000	NaN




Week 2018-10-29 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.848489	0.979026	0.909095	18642.000000	0.979026
1	0.490885	0.103685	0.171208	3636.000000	NaN
accuracy	0.836161	0.836161	0.836161	0.836161	0.836161
macro avg	0.669687	0.541356	0.540152	22278.000000	NaN
weighted avg	0.790124	0.836161	0.788665	22278.000000	NaN




Week 2018-11-05 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.857515	0.979933	0.914646	18787.000000	0.979933
1	0.533993	0.123747	0.200930	3491.000000	NaN
accuracy	0.845767	0.845767	0.845767	0.845767	0.845767
macro avg	0.695754	0.551840	0.557788	22278.000000	NaN
weighted avg	0.806819	0.845767	0.802806	22278.000000	NaN




Week 2018-11-12 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.861227	0.977728	0.915787	18858.000000	0.977728
1	0.516686	0.131287	0.209373	3420.000000	NaN
accuracy	0.847787	0.847787	0.847787	0.847787	0.847787
macro avg	0.688956	0.554507	0.562580	22278.000000	NaN
weighted avg	0.808335	0.847787	0.807342	22278.000000	NaN





Week 2018-11-19 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.852409	0.981329	0.912338	18639.000000	0.981329
1	0.575610	0.129706	0.211707	3639.000000	NaN
accuracy	0.842221	0.842221	0.842221	0.842221	0.842221
macro avg	0.714010	0.555518	0.562022	22278.000000	NaN
weighted avg	0.807196	0.842221	0.797893	22278.000000	NaN



Week 2018-11-26 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.860471	0.978426	0.915665	18726.000000	0.978426
1	0.589848	0.163570	0.256116	3552.000000	NaN
accuracy	0.848505	0.848505	0.848505	0.848505	0.848505
macro avg	0.725159	0.570998	0.585891	22278.000000	NaN
weighted avg	0.817323	0.848505	0.810507	22278.000000	NaN




```
[250]: # Fitting a Decision Tree
model = DecisionTreeClassifier(random_state=42, class_weight='balanced')
features = ['prev_week_activity',
            'prev_week_session_duration',
            'prev_dummy_level',
            'rolling_diff_1_weeks',
            'rolling_diff_2_weeks',
            'rolling_diff_3_weeks',
            'rolling_diff_4_weeks',
            'rolling_diff_5_weeks',
            'rolling_diff_6_weeks',
            'rolling_diff_7_weeks',
            'rolling_diff_8_weeks',
            'prev_week_rolling_diff_7_days',
            'prev_week_rolling_diff_14_days',
            'prev_week_rolling_diff_21_days',
            'prev_week_rolling_diff_28_days'
          ]
label = 'churn'

rolling_churn_prediction(df = df_all_weeks, features=features, label=label,
↳model=model)
```

Model type: <class 'sklearn.tree._classes.DecisionTreeClassifier'>


Week 2018-10-08 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.822201	1.000000	0.902426	18317.000000	1.000000
1	0.000000	0.000000	0.000000	3961.000000	NaN
accuracy	0.822201	0.822201	0.822201	0.822201	0.822201
macro avg	0.411101	0.500000	0.451213	22278.000000	NaN
weighted avg	0.676015	0.822201	0.741976	22278.000000	NaN




Week 2018-10-15 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.848715	0.925896	0.885627	18474.000000	0.925896
1	0.355461	0.198475	0.254723	3804.000000	NaN
accuracy	0.801688	0.801688	0.801688	0.801688	0.801688
macro avg	0.602088	0.562186	0.570175	22278.000000	NaN
weighted avg	0.764491	0.801688	0.777899	22278.000000	NaN




Week 2018-10-22 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.860744	0.896774	0.878389	18658.000000	0.896774
1	0.321592	0.252210	0.282706	3620.000000	NaN
accuracy	0.792037	0.792037	0.792037	0.792037	0.792037
macro avg	0.591168	0.574492	0.580548	22278.000000	NaN
weighted avg	0.773136	0.792037	0.781596	22278.000000	NaN




Week 2018-10-29 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.864050	0.886761	0.875258	18642.000000	0.886761
1	0.328989	0.284653	0.305220	3636.000000	NaN
accuracy	0.788491	0.788491	0.788491	0.788491	0.788491
macro avg	0.596519	0.585707	0.590239	22278.000000	NaN
weighted avg	0.776722	0.788491	0.782222	22278.000000	NaN




Week 2018-11-05 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.873463	0.896152	0.884662	18787.000000	0.896152
1	0.350316	0.301346	0.323991	3491.000000	NaN
accuracy	0.802945	0.802945	0.802945	0.802945	0.802945
macro avg	0.611890	0.598749	0.604327	22278.000000	NaN
weighted avg	0.791485	0.802945	0.796804	22278.000000	NaN




Week 2018-11-12 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.878913	0.891823	0.885321	18858.000000	0.891823
1	0.350939	0.322515	0.336127	3420.000000	NaN
accuracy	0.804426	0.804426	0.804426	0.804426	0.804426
macro avg	0.614926	0.607169	0.610724	22278.000000	NaN
weighted avg	0.797861	0.804426	0.801012	22278.000000	NaN





Week 2018-11-19 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.872636	0.896185	0.884254	18639.000000	0.896185
1	0.382972	0.330036	0.354539	3639.000000	NaN
accuracy	0.803708	0.803708	0.803708	0.803708	0.803708
macro avg	0.627804	0.613111	0.619396	22278.000000	NaN
weighted avg	0.792652	0.803708	0.797728	22278.000000	NaN



Week 2018-11-26 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.879556	0.896935	0.888160	18726.000000	0.896935
1	0.393463	0.352477	0.371844	3552.000000	NaN
accuracy	0.810127	0.810127	0.810127	0.810127	0.810127
macro avg	0.636510	0.624706	0.630002	22278.000000	NaN
weighted avg	0.802053	0.810127	0.805839	22278.000000	NaN



```
[ ]: print(df_all_weeks[features + [label]].corr()[label].
      ↪sort_values(ascending=False))
```

```
churn                1.000000
prev_week_session_duration  0.273998
rolling_diff_4_weeks    0.156740
rolling_diff_5_weeks    0.155946
rolling_diff_3_weeks    0.155399
rolling_diff_2_weeks    0.153361
rolling_diff_6_weeks    0.152086
rolling_diff_7_weeks    0.149336
rolling_diff_8_weeks    0.147695
prev_week_activity      0.142876
rolling_diff_1_weeks    0.138172
prev_dummy_level       -0.037114
Name: churn, dtype: float64
```


```
[251]: model = MLPClassifier(hidden_layer_sizes=(100, 50), max_iter=500,
      ↪random_state=42)
features = ['prev_week_activity',
            'prev_week_session_duration',
            'prev_dummy_level',
            'rolling_diff_1_weeks',
            'rolling_diff_2_weeks',
            'rolling_diff_3_weeks',
            'rolling_diff_4_weeks',
            'rolling_diff_5_weeks',
            'rolling_diff_6_weeks',
            'rolling_diff_7_weeks',
            'rolling_diff_8_weeks',
            'prev_week_rolling_diff_7_days',
            'prev_week_rolling_diff_14_days',
            'prev_week_rolling_diff_21_days',
            'prev_week_rolling_diff_28_days'
            ]
label = 'churn'

rolling_churn_prediction(df = df_all_weeks, features=features, label=label,
      ↪model=model)
```

```
Model type: <class
'sklearn.neural_network._multilayer_perceptron.MLPClassifier'>
Week 2018-10-08 00:00:00:
```


	precision	recall	f1-score	support	specificity
0	0.822201	1.000000	0.902426	18317.000000	1.000000
1	0.000000	0.000000	0.000000	3961.000000	NaN

accuracy	0.822201	0.822201	0.822201	0.822201	0.822201
macro avg	0.411101	0.500000	0.451213	22278.000000	NaN
weighted avg	0.676015	0.822201	0.741976	22278.000000	NaN




Week 2018-10-15 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.829479	0.999783	0.906703	18474.000000	0.999783
1	0.636364	0.001840	0.003670	3804.000000	NaN
accuracy	0.829383	0.829383	0.829383	0.829383	0.829383
macro avg	0.732921	0.500812	0.455187	22278.000000	NaN
weighted avg	0.796504	0.829383	0.752509	22278.000000	NaN




Week 2018-10-22 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.838390	0.998178	0.911333	18658.000000	0.998178
1	0.468750	0.008287	0.016287	3620.000000	NaN
accuracy	0.837328	0.837328	0.837328	0.837328	0.837328
macro avg	0.653570	0.503233	0.463810	22278.000000	NaN
weighted avg	0.778327	0.837328	0.765895	22278.000000	NaN




Week 2018-10-29 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.845080	0.990505	0.912032	18642.000000	0.990505
1	0.586449	0.069032	0.123524	3636.000000	NaN
accuracy	0.840111	0.840111	0.840111	0.840111	0.840111
macro avg	0.715764	0.529769	0.517778	22278.000000	NaN
weighted avg	0.802869	0.840111	0.783339	22278.000000	NaN




Week 2018-11-05 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.992777	0.475542	0.643058	18787.000000	0.475542
1	0.258001	0.981381	0.408587	3491.000000	NaN
accuracy	0.554807	0.554807	0.554807	0.554807	0.554807
macro avg	0.625389	0.728461	0.525822	22278.000000	NaN
weighted avg	0.877636	0.554807	0.606316	22278.000000	NaN




Week 2018-11-12 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.841679	0.942995	0.889461	18858.000000	0.942995
1	0.065217	0.021930	0.032823	3420.000000	NaN
accuracy	0.801598	0.801598	0.801598	0.801598	0.801598
macro avg	0.453448	0.482462	0.461142	22278.000000	NaN
weighted avg	0.722481	0.801598	0.757955	22278.000000	NaN



Week 2018-11-19 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.835700	0.940930	0.885199	18639.000000	0.940930
1	0.147833	0.052487	0.077469	3639.000000	NaN
accuracy	0.795808	0.795808	0.795808	0.795808	0.795808
macro avg	0.491766	0.496709	0.481334	22278.000000	NaN
weighted avg	0.723340	0.795808	0.753260	22278.000000	NaN



Week 2018-11-26 00:00:00:

	precision	recall	f1-score	support	specificity
0	0.841443	0.999252	0.913583	18726.000000	0.999252
1	0.650000	0.007320	0.014477	3552.000000	NaN
accuracy	0.841099	0.841099	0.841099	0.841099	0.841099
macro avg	0.745721	0.503286	0.464030	22278.000000	NaN
weighted avg	0.810919	0.841099	0.770229	22278.000000	NaN

